

Notebook

October 11, 2024

1 HW1

Author: QIhua Dong

Github: <https://github.com/dddraxxx/ee5644>

2 Q1

2.0.1 Generate a Gaussian mixture dataset

```
[2]: import numpy as np
from scipy.stats import multivariate_normal

# Parameters for class L=0
m0 = np.array([-1, -1, -1, -1])
C0 = np.array([[2, -0.5, 0.3, 0],
               [-0.5, 1, -0.5, 0],
               [0.3, -0.5, 1, 0],
               [0, 0, 0, 2]])

# Parameters for class L=1
m1 = np.array([1, 1, 1, 1])
C1 = np.array([[1, 0.3, -0.2, 0],
               [0.3, 2, 0.3, 0],
               [-0.2, 0.3, 1, 0],
               [0, 0, 0, 3]])

# Class priors
P_L0 = 0.35
P_L1 = 0.65

# Number of samples
num_samples = 10000

# Generate class labels L (0 or 1) based on the priors
labels = np.random.choice([0, 1], size=num_samples, p=[P_L0, P_L1])

# Initialize an empty array to store the generated samples
```

```

samples = np.zeros((num_samples, 4))

# Generate samples for each class based on the Gaussian distributions
for i in range(num_samples):
    if labels[i] == 0:
        samples[i] = multivariate_normal.rvs(mean=m0, cov=C0)
    else:
        samples[i] = multivariate_normal.rvs(mean=m1, cov=C1)

# Save the generated samples and labels to a file
np.savez('gaussian_mixture_samples.npz', samples=samples, labels=labels)

print(f"Generated {num_samples} samples and saved to 'gaussian_mixture_samples.
↪npz'")

```

Generated 10000 samples and saved to 'gaussian_mixture_samples.npz'

```
[4]: samples, labels
```

```

[4]: (array([[ 1.21702054,  0.42134373,  0.10965336,  1.14882005],
 [ 3.86235031,  4.40286043,  2.04268341,  2.08843631],
 [-0.39410192, -0.41088299, -2.98079201, -1.47109445],
 ...,
 [-0.69416581, -2.65615592,  1.06310199, -3.30495517],
 [ 0.69585596, -1.63946255,  0.50437631,  2.381229  ],
 [ 0.26764186,  4.44565221,  2.03894145,  0.79276422]]),
 array([1, 1, 0, ..., 1, 1, 1]))

```

2.0.2 Part A

```

[12]: # Import necessary libraries
import numpy as np
from scipy.stats import multivariate_normal

# Load the generated data from part 1
data = np.load('gaussian_mixture_samples.npz')
samples = data['samples']
labels = data['labels']

# Compute the likelihoods  $p(x/L=0)$  and  $p(x/L=1)$  for each sample
p_x_given_L0 = multivariate_normal.pdf(samples, mean=m0, cov=C0)
p_x_given_L1 = multivariate_normal.pdf(samples, mean=m1, cov=C1)

# Likelihood ratio for each sample
likelihood_ratio = p_x_given_L1 / p_x_given_L0

# Print the answer for step 1
print("Minimum Expected Risk Classification Rule")

```

```
print("Likelihood ratio computed as  $p(x|L=1) / p(x|L=0)$  for each sample.")
print(f"Sample likelihood ratios: {likelihood_ratio}") # Print the first 5
↳ likelihood ratios as a sample
```

Minimum Expected Risk Classification Rule

Likelihood ratio computed as $p(x|L=1) / p(x|L=0)$ for each sample.

Sample likelihood ratios: [3.51761361e+02 1.89713759e+20 7.80182270e-05 ...
2.50332324e-02
7.99132323e+00 1.64115809e+16]

```
[16]: import numpy as np
import matplotlib.pyplot as plt

# Set up a range of gamma (threshold) values to sweep through
gamma_values = np.logspace(-3, 3, num=500) # 500 gamma values from  $10^{-3}$  to
↳  $10^3$ 

# Lists to store true positive and false positive rates for the ROC curve
tpr_values = [] # True Positive Rate ( $P(D=1 | L=1)$ )
fpr_values = [] # False Positive Rate ( $P(D=1 | L=0)$ )

# Iterate through each gamma and compute TPR and FPR
for gamma in gamma_values:
    # Apply the likelihood ratio test: decide class based on threshold gamma
    decisions = (likelihood_ratio > gamma).astype(int)

    # True positives:  $D=1$  and  $L=1$ 
    tp = np.sum((decisions == 1) & (labels == 1))
    fn = np.sum((decisions == 0) & (labels == 1))
    tpr = tp / (tp + fn) # True positive rate

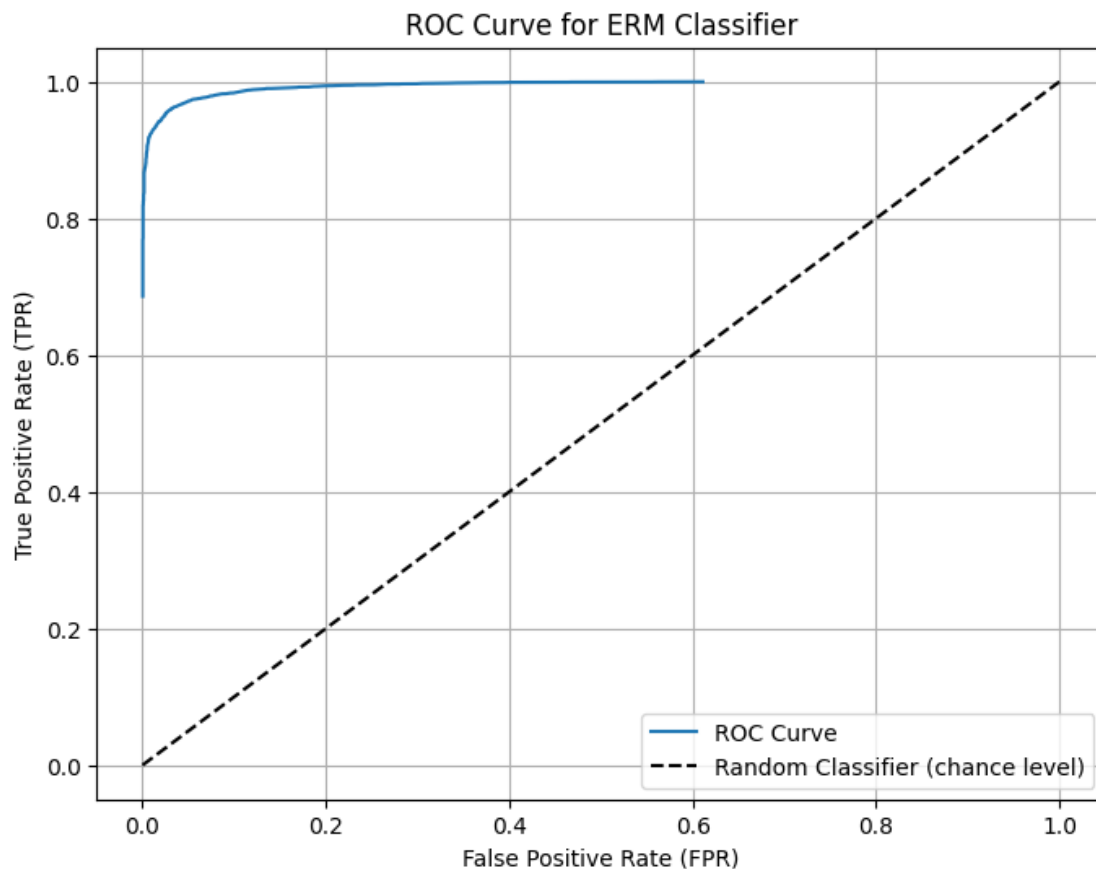
    # False positives:  $D=1$  and  $L=0$ 
    fp = np.sum((decisions == 1) & (labels == 0))
    tn = np.sum((decisions == 0) & (labels == 0))
    fpr = fp / (fp + tn) # False positive rate

    tpr_values.append(tpr)
    fpr_values.append(fpr)

# Plot the ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr_values, tpr_values, label='ROC Curve')
plt.plot([0, 1], [0, 1], 'k--', label='Random Classifier (chance level)')
plt.title('ROC Curve for ERM Classifier')
plt.xlabel('False Positive Rate (FPR)')
plt.ylabel('True Positive Rate (TPR)')
plt.legend(loc='lower right')
```

```
plt.grid()
plt.show()

# Print the answer for step 2
print("Step 2: ROC Curve computed and plotted.")
# print(f"TPR values (sample): {np.array(tpr_values)}")
# print(f"FPR values (sample): {np.array(fpr_values)}")
```



Step 2: ROC Curve computed and plotted.

```
[19]: # Step 3: Find the gamma that minimizes the probability of error
#  $P(\text{error}; \gamma) = P(D=1 \mid L=0) * P(L=0) + P(D=0 \mid L=1) * P(L=1)$ 
P_L0 = 0.35
P_L1 = 0.65

# Initialize an empty list to store the probability of error for each gamma
errors = []

# Calculate the error for each gamma value
```

```

for i, gamma in enumerate(gamma_values):
    # False positive rate and false negative rate (1 - True positive rate)
    fpr = fpr_values[i]
    fnr = 1 - tpr_values[i]

    # Probability of error for this gamma
    p_error = fpr * P_L0 + fnr * P_L1
    errors.append(p_error)

# Find the index of the minimum error
min_error_idx = np.argmin(errors)
min_error_gamma = gamma_values[min_error_idx]
min_error_value = errors[min_error_idx]

# Plot the ROC curve again, but highlight the point of minimum error
plt.figure(figsize=(8, 6))
plt.plot(fpr_values, tpr_values, label='ROC Curve')
plt.plot([0, 1], [0, 1], 'k--', label='Random Classifier (chance level)')
plt.scatter(fpr_values[min_error_idx], tpr_values[min_error_idx], color='red',
            label=f'Min Error (gamma={min_error_gamma:.2f})', zorder=5)
plt.title('ROC Curve with Min Error Point Highlighted')
plt.xlabel('False Positive Rate (FPR)')
plt.ylabel('True Positive Rate (TPR)')
plt.legend(loc='lower right')
plt.grid()
plt.show()

# Print the answer for step 3
print("Step 3: Minimizing the Probability of Error")
print(f"The minimum probability of error is {min_error_value:.4f} at gamma = {min_error_gamma:.4f}")
print(f"Corresponding FPR: {fpr_values[min_error_idx]:.4f}, TPR: {tpr_values[min_error_idx]:.4f}")

# Compute the theoretical optimal gamma
gamma_opt_theoretical = P_L0 / P_L1

# Print the comparison
print("\n")
print("Comparison of Empirically Selected Gamma and Theoretical Optimal Gamma")
print(f"Empirically selected gamma (from ROC curve) = {min_error_gamma:.4f}")
print(f"Theoretically optimal gamma (from priors) = {gamma_opt_theoretical:.4f}")

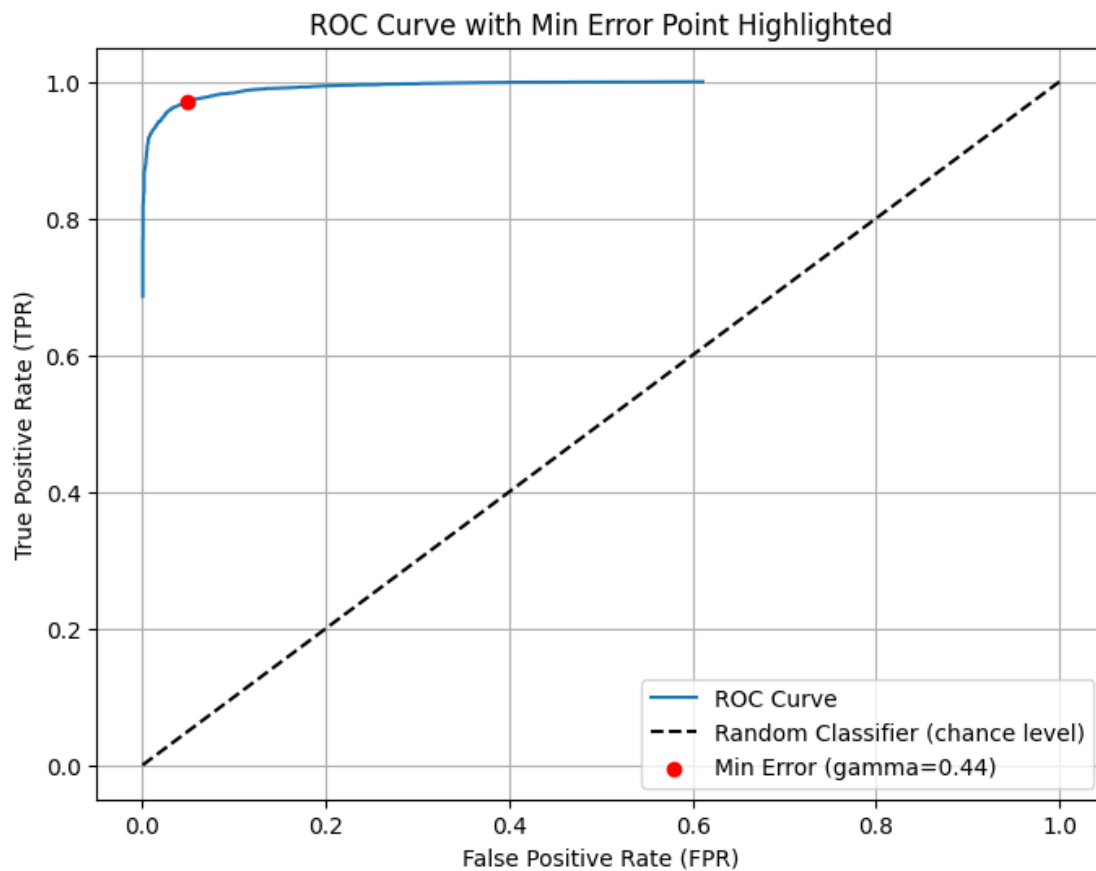
# Compare the results
if np.isclose(min_error_gamma, gamma_opt_theoretical, atol=0.01):
    print("The empirical gamma is very close to the theoretical optimal gamma.")

```

```

else:
    print("The empirical gamma differs from the theoretical optimal gamma.")

```



Step 3: Minimizing the Probability of Error

The minimum probability of error is 0.0358 at $\gamma = 0.4419$

Corresponding FPR: 0.0494, TPR: 0.9715

Comparison of Empirically Selected Gamma and Theoretical Optimal Gamma

Empirically selected gamma (from ROC curve) = 0.4419

Theoretically optimal gamma (from priors) = 0.5385

The empirical gamma differs from the theoretical optimal gamma.

2.0.3 Part B

```

[20]: import numpy as np
      from scipy.stats import multivariate_normal
      import matplotlib.pyplot as plt

      # Load the generated data

```

```

data = np.load('gaussian_mixture_samples.npz')
samples = data['samples']
labels = data['labels']

# True mean vectors from Part A
m0 = np.array([-1, -1, -1, -1])
m1 = np.array([1, 1, 1, 1])

# True covariance matrices from Part A
C0_true = np.array([[2, -0.5, 0.3, 0],
                    [-0.5, 1, -0.5, 0],
                    [0.3, -0.5, 1, 0],
                    [0, 0, 0, 2]])

C1_true = np.array([[1, 0.3, -0.2, 0],
                    [0.3, 2, 0.3, 0],
                    [-0.2, 0.3, 1, 0],
                    [0, 0, 0, 3]])

# Naive Bayesian assumption: use diagonal covariance matrices
# Extract the diagonal entries (variances) for each class
C0_naive = np.diag(np.diag(C0_true)) # Diagonal matrix for class L=0
C1_naive = np.diag(np.diag(C1_true)) # Diagonal matrix for class L=1

# Print the diagonal covariance matrices to verify
print("Diagonal covariance matrix for class L=0 (Naive Bayes assumption):")
print(C0_naive)
print("Diagonal covariance matrix for class L=1 (Naive Bayes assumption):")
print(C1_naive)

```

Diagonal covariance matrix for class L=0 (Naive Bayes assumption):

```

[[2. 0. 0. 0.]
 [0. 1. 0. 0.]
 [0. 0. 1. 0.]
 [0. 0. 0. 2.]]

```

Diagonal covariance matrix for class L=1 (Naive Bayes assumption):

```

[[1. 0. 0. 0.]
 [0. 2. 0. 0.]
 [0. 0. 1. 0.]
 [0. 0. 0. 3.]]

```

```

[22]: # Class priors
P_L0 = 0.35
P_L1 = 0.65

# Compute the likelihoods  $p(x/L=0)$  and  $p(x/L=1)$  using the Naive Bayes assumption
p_x_given_L0_naive = multivariate_normal.pdf(samples, mean=m0, cov=C0_naive)

```

```

p_x_given_L1_naive = multivariate_normal.pdf(samples, mean=m1, cov=C1_naive)

# Likelihood ratio for each sample under Naive Bayes assumption
likelihood_ratio_naive = p_x_given_L1_naive / p_x_given_L0_naive

# Print some sample likelihood ratios
print(f"Sample likelihood ratios (Naive Bayes assumption):␣
↪{likelihood_ratio_naive}")

```

```

Sample likelihood ratios (Naive Bayes assumption): [2.70824011e+01
3.20203944e+08 2.41014784e-04 ... 3.99848265e-02
1.19619104e+01 1.72700163e+07]

```

```

[23]: # Set up a range of gamma (threshold) values to sweep through
gamma_values = np.logspace(-3, 3, num=500) # 500 gamma values from 10-3 to␣
↪103

# Lists to store true positive and false positive rates for the ROC curve
tpr_values_naive = [] # True Positive Rate (P(D=1 | L=1))
fpr_values_naive = [] # False Positive Rate (P(D=1 | L=0))

# Iterate through each gamma and compute TPR and FPR using Naive Bayes␣
↪assumption
for gamma in gamma_values:
    decisions_naive = (likelihood_ratio_naive > gamma).astype(int)

    # True positives: D=1 and L=1
    tp = np.sum((decisions_naive == 1) & (labels == 1))
    fn = np.sum((decisions_naive == 0) & (labels == 1))
    tpr_naive = tp / (tp + fn) # True positive rate

    # False positives: D=1 and L=0
    fp = np.sum((decisions_naive == 1) & (labels == 0))
    tn = np.sum((decisions_naive == 0) & (labels == 0))
    fpr_naive = fp / (fp + tn) # False positive rate

    tpr_values_naive.append(tpr_naive)
    fpr_values_naive.append(fpr_naive)

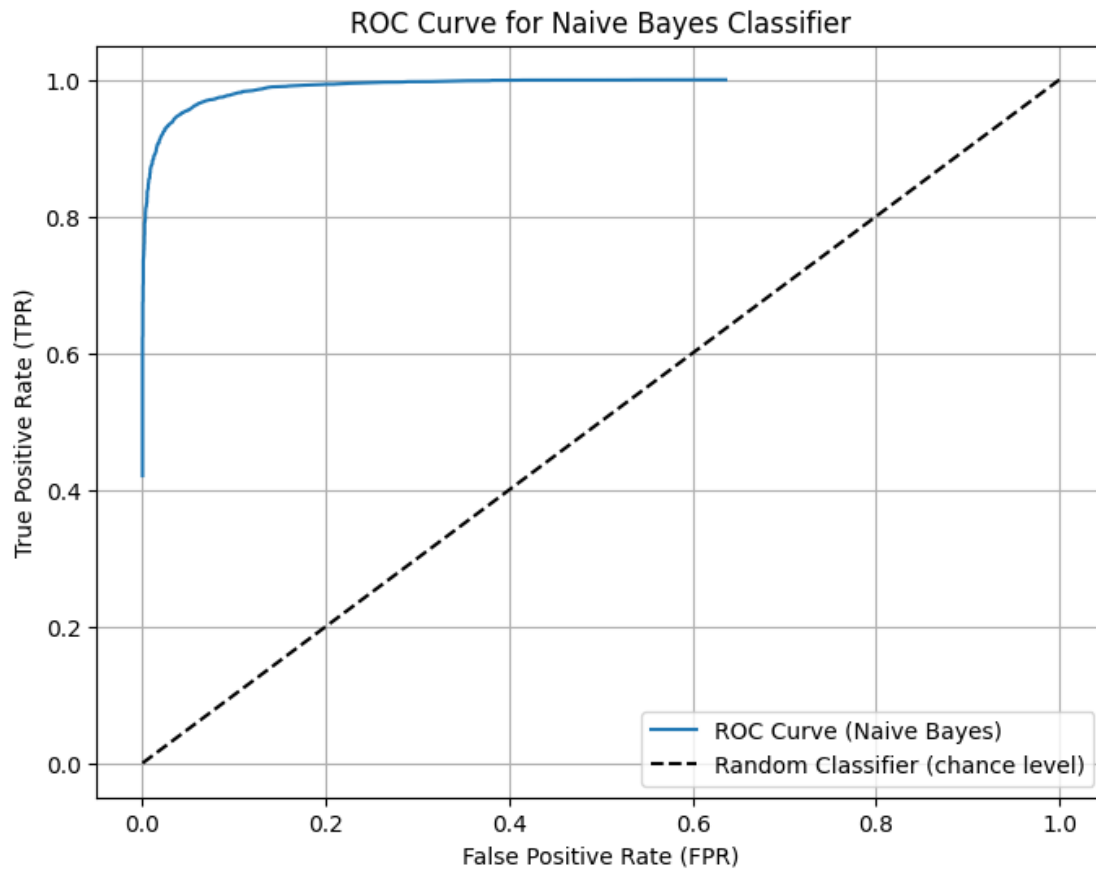
# Plot the ROC curve for the Naive Bayes assumption
plt.figure(figsize=(8, 6))
plt.plot(fpr_values_naive, tpr_values_naive, label='ROC Curve (Naive Bayes)')
plt.plot([0, 1], [0, 1], 'k--', label='Random Classifier (chance level)')
plt.title('ROC Curve for Naive Bayes Classifier')
plt.xlabel('False Positive Rate (FPR)')
plt.ylabel('True Positive Rate (TPR)')
plt.legend(loc='lower right')

```



```
plt.grid()
plt.show()

# Print the ROC curve details
print("ROC Curve for Naive Bayes classifier plotted.")
```



ROC Curve for Naive Bayes classifier plotted.

```
[24]: # Step 4: Find the gamma that minimizes the probability of error using Naive
      ↪ Bayes assumption
errors_naive = []

# Calculate the error for each gamma value using Naive Bayes assumption
for i, gamma in enumerate(gamma_values):
    # False positive rate and false negative rate (1 - True positive rate)
    fpr_naive = fpr_values_naive[i]
    fnr_naive = 1 - tpr_values_naive[i]

    # Probability of error for this gamma
```

```

p_error_naive = fpr_naive * P_L0 + fnr_naive * P_L1
errors_naive.append(p_error_naive)

# Find the index of the minimum error
min_error_idx_naive = np.argmin(errors_naive)
min_error_gamma_naive = gamma_values[min_error_idx_naive]
min_error_value_naive = errors_naive[min_error_idx_naive]

# Print the results
print("Step 4: Minimizing the Probability of Error with Naive Bayes Assumption")
print(f"The minimum probability of error (Naive Bayes) is_
↳{min_error_value_naive:.4f} at gamma = {min_error_gamma_naive:.4f}")

```

Step 4: Minimizing the Probability of Error with Naive Bayes Assumption
The minimum probability of error (Naive Bayes) is 0.0438 at gamma = 0.4298

```

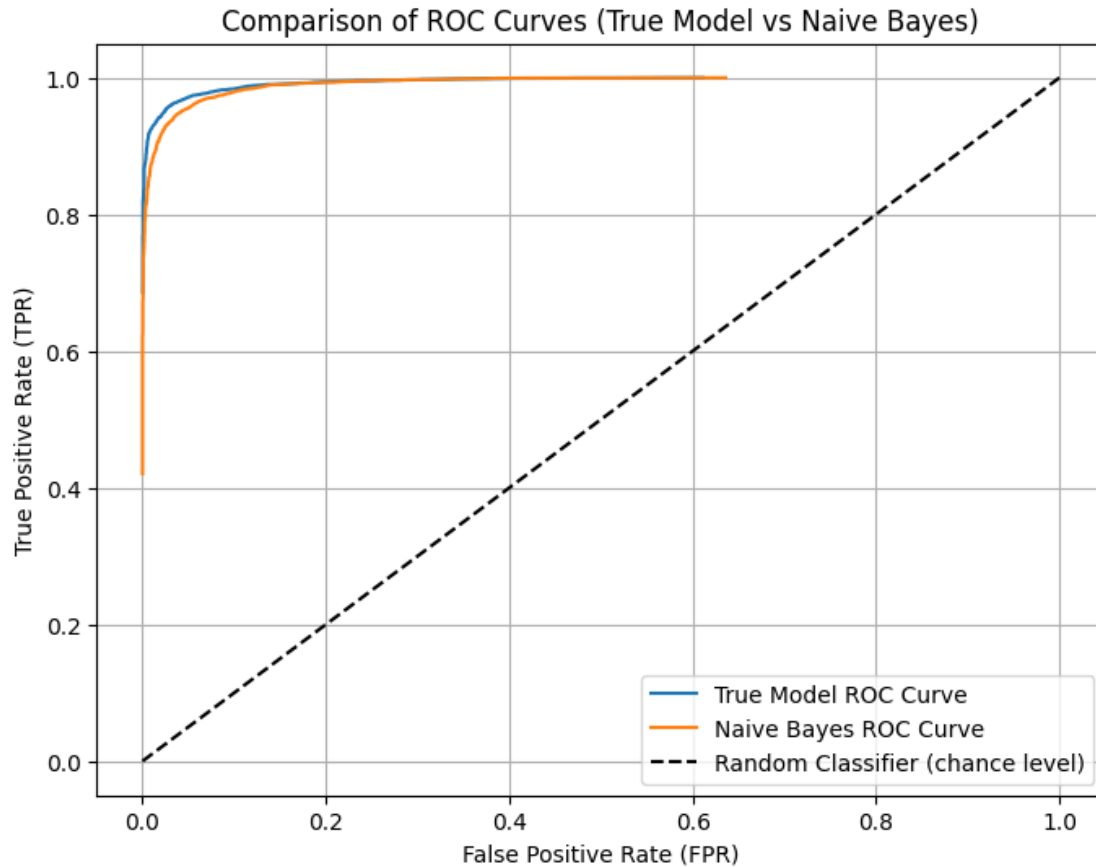
[25]: # Plot ROC Curves for both models to compare
plt.figure(figsize=(8, 6))
plt.plot(fpr_values, tpr_values, label='True Model ROC Curve')
plt.plot(fpr_values_naive, tpr_values_naive, label='Naive Bayes ROC Curve')
plt.plot([0, 1], [0, 1], 'k--', label='Random Classifier (chance level)')
plt.title('Comparison of ROC Curves (True Model vs Naive Bayes)')
plt.xlabel('False Positive Rate (FPR)')
plt.ylabel('True Positive Rate (TPR)')
plt.legend(loc='lower right')
plt.grid()
plt.show()

# Print minimum error comparison
print("Comparison of Minimum Probability of Error")
print(f"True Model: Minimum probability of error = {min(errors):.4f} at gamma =_
↳{gamma_values[np.argmin(errors)]:.4f}")
print(f"Naive Bayes Model: Minimum probability of error =_
↳{min_error_value_naive:.4f} at gamma = {min_error_gamma_naive:.4f}")

# Analyze if the error rates differ significantly
error_diff = abs(min(errors) - min_error_value_naive)
print(f"Difference in minimum error between True Model and Naive Bayes:_
↳{error_diff:.4f}")

# Conclusion
if error_diff > 0.01:
    print("The Naive Bayes assumption negatively impacted the performance,_
↳leading to a higher probability of error.")
else:
    print("The Naive Bayes assumption had a minimal impact on the performance,_
↳in this case.")

```



Comparison of Minimum Probability of Error

True Model: Minimum probability of error = 0.0358 at $\gamma = 0.4419$

Naive Bayes Model: Minimum probability of error = 0.0438 at $\gamma = 0.4298$

Difference in minimum error between True Model and Naive Bayes: 0.0080

The Naive Bayes assumption had a minimal impact on the performance in this case.

2.1 Part C

```
[26]: import numpy as np

# Load the generated data
data = np.load('gaussian_mixture_samples.npz')
samples = data['samples']
labels = data['labels']

# Separate the samples by class
class_0_samples = samples[labels == 0]
class_1_samples = samples[labels == 1]

# Estimate the means (sample average for each class)
```

```

mean_0 = np.mean(class_0_samples, axis=0)
mean_1 = np.mean(class_1_samples, axis=0)

# Estimate the covariances (sample covariance for each class)
cov_0 = np.cov(class_0_samples, rowvar=False)
cov_1 = np.cov(class_1_samples, rowvar=False)

# Print the estimated means and covariances
print("Estimated mean for class 0:", mean_0)
print("Estimated mean for class 1:", mean_1)
print("Estimated covariance for class 0:\n", cov_0)
print("Estimated covariance for class 1:\n", cov_1)

```

Estimated mean for class 0: [-0.97885076 -1.0237032 -0.9936786 -0.98522708]

Estimated mean for class 1: [0.98877577 1.01837441 1.00133738 1.03811884]

Estimated covariance for class 0:

```

[[ 2.05628253 -0.50954691  0.30881279  0.02713176]
 [-0.50954691  0.99377434 -0.50324397 -0.0270095 ]
 [ 0.30881279 -0.50324397  1.01671457  0.02585526]
 [ 0.02713176 -0.0270095   0.02585526  1.93776272]]

```

Estimated covariance for class 1:

```

[[ 0.98596894  0.28232123 -0.17598996 -0.00444704]
 [ 0.28232123  2.00259907  0.33323045 -0.02692974]
 [-0.17598996  0.33323045  1.00285683  0.03258915]
 [-0.00444704 -0.02692974  0.03258915  2.96640964]]

```

```

[27]: # Compute the within-class scatter matrix (S_W)
S_W = cov_0 + cov_1

# Compute the between-class difference in means
mean_diff = mean_1 - mean_0

# Compute the Fisher LDA weight vector
w_LDA = np.linalg.inv(S_W).dot(mean_diff)

# Print the Fisher LDA weight vector
print("Fisher LDA projection vector (w_LDA):", w_LDA)

```

Fisher LDA projection vector (w_LDA): [0.65951191 0.79556628 0.99968001 0.40636247]

```

[29]: # Project all the samples onto the LDA direction
projected_data = samples.dot(w_LDA)

# Print a few sample projections
print(f"First 5 projected data points: {projected_data}")

```

First 5 projected data points: [1.71430203 8.94072525 -4.16443533 ...
-2.85120666 0.62648119

6.07376262]

```
[30]: import matplotlib.pyplot as plt

# Set up a range of tau (threshold) values to sweep through
tau_values = np.linspace(np.min(projected_data), np.max(projected_data),
    ↪ num=500)

# Lists to store true positive and false positive rates for the ROC curve
tpr_values_lda = [] # True Positive Rate (P(D=1 | L=1))
fpr_values_lda = [] # False Positive Rate (P(D=1 | L=0))

# Iterate through each tau and compute TPR and FPR
for tau in tau_values:
    decisions_lda = (projected_data > tau).astype(int)

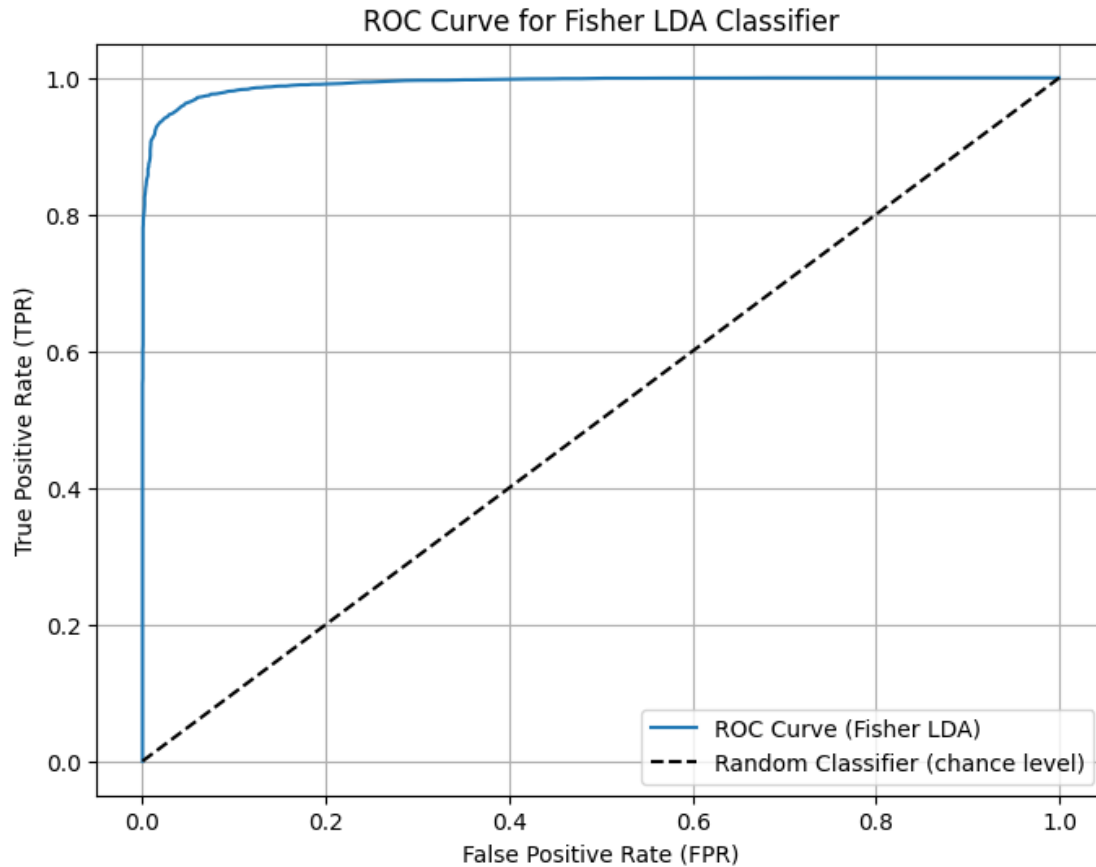
    # True positives: D=1 and L=1
    tp = np.sum((decisions_lda == 1) & (labels == 1))
    fn = np.sum((decisions_lda == 0) & (labels == 1))
    tpr_lda = tp / (tp + fn) # True positive rate

    # False positives: D=1 and L=0
    fp = np.sum((decisions_lda == 1) & (labels == 0))
    tn = np.sum((decisions_lda == 0) & (labels == 0))
    fpr_lda = fp / (fp + tn) # False positive rate

    tpr_values_lda.append(tpr_lda)
    fpr_values_lda.append(fpr_lda)

# Plot the ROC curve for Fisher LDA
plt.figure(figsize=(8, 6))
plt.plot(fpr_values_lda, tpr_values_lda, label='ROC Curve (Fisher LDA)')
plt.plot([0, 1], [0, 1], 'k--', label='Random Classifier (chance level)')
plt.title('ROC Curve for Fisher LDA Classifier')
plt.xlabel('False Positive Rate (FPR)')
plt.ylabel('True Positive Rate (TPR)')
plt.legend(loc='lower right')
plt.grid()
plt.show()

# Print the ROC curve details
print("ROC Curve for Fisher LDA classifier plotted.")
```



ROC Curve for Fisher LDA classifier plotted.

```
[31]: # Class priors
P_L0 = 0.35
P_L1 = 0.65

# Initialize an empty list to store the probability of error for each tau
errors_lda = []

# Calculate the error for each tau value
for i, tau in enumerate(tau_values):
    # False positive rate and false negative rate (1 - True positive rate)
    fpr_lda = fpr_values_lda[i]
    fnr_lda = 1 - tpr_values_lda[i]

    # Probability of error for this tau
    p_error_lda = fpr_lda * P_L0 + fnr_lda * P_L1
    errors_lda.append(p_error_lda)

# Find the index of the minimum error
```

```

min_error_idx_lda = np.argmin(errors_lda)
min_error_tau_lda = tau_values[min_error_idx_lda]
min_error_value_lda = errors_lda[min_error_idx_lda]

# Print the results
print("Step 5: Minimizing the Probability of Error with Fisher LDA")
print(f"The minimum probability of error (LDA) is {min_error_value_lda:.4f} at_\n
↳tau = {min_error_tau_lda:.4f}")

```

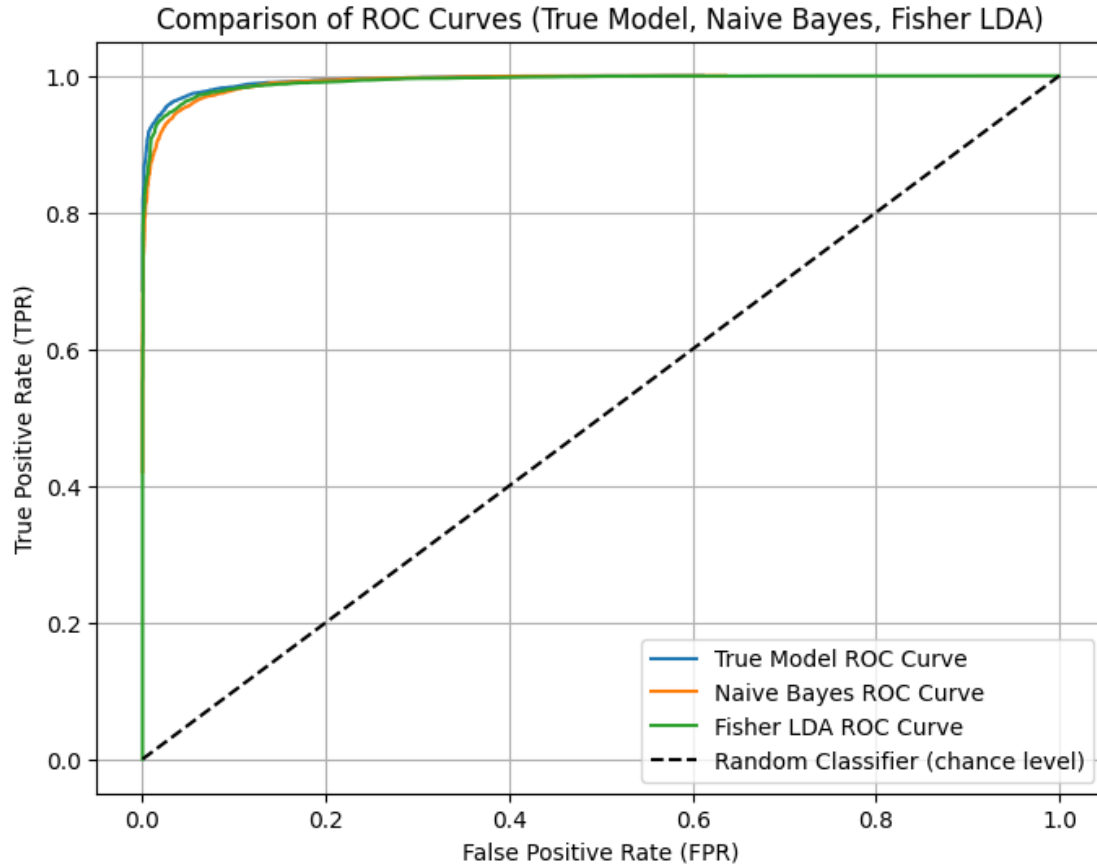
Step 5: Minimizing the Probability of Error with Fisher LDA
The minimum probability of error (LDA) is 0.0397 at tau = -0.7305

```

[33]: # Plot ROC Curves for all three models (True Model, Naive Bayes, Fisher LDA)_\n
↳for comparison
plt.figure(figsize=(8, 6))
plt.plot(fpr_values, tpr_values, label='True Model ROC Curve')
plt.plot(fpr_values_naive, tpr_values_naive, label='Naive Bayes ROC Curve')
plt.plot(fpr_values_lda, tpr_values_lda, label='Fisher LDA ROC Curve')
plt.plot([0, 1], [0, 1], 'k--', label='Random Classifier (chance level)')
plt.title('Comparison of ROC Curves (True Model, Naive Bayes, Fisher LDA)')
plt.xlabel('False Positive Rate (FPR)')
plt.ylabel('True Positive Rate (TPR)')
plt.legend(loc='lower right')
plt.grid()
plt.show()

# Print minimum error comparison for all three models
print("Comparison of Minimum Probability of Error")
print(f"True Model: Minimum probability of error = {min(errors):.4f} at gamma =_\n
↳{gamma_values[np.argmin(errors)]:.4f}")
print(f"Naive Bayes Model: Minimum probability of error =_\n
↳{min_error_value_naive:.4f} at gamma = {min_error_gamma_naive:.4f}")
print(f"Fisher LDA Model: Minimum probability of error = {min_error_value_lda:.\n
↳4f} at tau = {min_error_tau_lda:.4f}")

```



Comparison of Minimum Probability of Error

True Model: Minimum probability of error = 0.0358 at $\gamma = 0.4419$

Naive Bayes Model: Minimum probability of error = 0.0438 at $\gamma = 0.4298$

Fisher LDA Model: Minimum probability of error = 0.0397 at $\tau = -0.7305$

2.1.1 LDA Classifier Performance Compared to True Model and Naive Bayes

Summary:

- **True Model:** Optimal performance due to full knowledge of class distributions.
- **Naive Bayes:** Suffers from feature independence assumption.
- **Fisher LDA:** Provides a strong balance between simplicity and performance, offering a near-optimal solution while being computationally efficient.

Notebook

October 11, 2024

1 Q2

1.1 Part A: Minimum Probability of Error Classification (0-1 Loss)

```
[2]: ## Step 1: Generate 10,000 Samples from the Data Distribution
import numpy as np

# Class priors
P_L1 = 0.3
P_L2 = 0.3
P_L3 = 0.4

# Means and covariances for each class
mean_1 = [0, 0, 0] # Mean for class 1
mean_2 = [3, 3, 3] # Mean for class 2
mean_3a = [6, 0, 0] # First component of class 3
mean_3b = [0, 6, 6] # Second component of class 3

cov_1 = np.eye(3) # Covariance for class 1 (identity)
cov_2 = np.eye(3) # Covariance for class 2 (identity)
cov_3 = np.eye(3) # Same covariance for both components of class 3

# Number of samples to generate
n_samples = 10000

# Generate samples for each class based on the priors
n_L1 = int(P_L1 * n_samples)
n_L2 = int(P_L2 * n_samples)
n_L3 = n_samples - n_L1 - n_L2

# Generate class 1 samples (single Gaussian)
samples_L1 = np.random.multivariate_normal(mean_1, cov_1, n_L1)

# Generate class 2 samples (single Gaussian)
samples_L2 = np.random.multivariate_normal(mean_2, cov_2, n_L2)

# Generate class 3 samples (from mixture of two Gaussians)
samples_L3a = np.random.multivariate_normal(mean_3a, cov_3, n_L3 // 2)
```

```

samples_L3b = np.random.multivariate_normal(mean_3b, cov_3, n_L3 // 2)
samples_L3 = np.vstack((samples_L3a, samples_L3b))

# Combine samples and create true labels
samples = np.vstack((samples_L1, samples_L2, samples_L3))
labels = np.array([1] * n_L1 + [2] * n_L2 + [3] * n_L3)

# Print shapes to verify
print(f"Generated samples shape: {samples.shape}")
print(f"Generated labels shape: {labels.shape}")

```

Generated samples shape: (10000, 3)

Generated labels shape: (10000,)

```

[3]: ## Step 2: Implement the Bayesian Decision Rule (Minimum Error)
from scipy.stats import multivariate_normal

# Compute the class-conditional likelihoods for each class
likelihood_L1 = multivariate_normal.pdf(samples, mean=mean_1, cov=cov_1)
likelihood_L2 = multivariate_normal.pdf(samples, mean=mean_2, cov=cov_2)
likelihood_L3a = multivariate_normal.pdf(samples, mean=mean_3a, cov=cov_3)
likelihood_L3b = multivariate_normal.pdf(samples, mean=mean_3b, cov=cov_3)

# Class 3 is a mixture, so we average the likelihoods from the two components
likelihood_L3 = 0.5 * (likelihood_L3a + likelihood_L3b)

# Compute the posterior probabilities using Bayes' Rule
posterior_L1 = likelihood_L1 * P_L1
posterior_L2 = likelihood_L2 * P_L2
posterior_L3 = likelihood_L3 * P_L3

# Combine posteriors into a matrix
posteriors = np.vstack((posterior_L1, posterior_L2, posterior_L3)).T

# Classify based on maximum posterior probability (Bayesian Decision Rule)
predicted_labels = np.argmax(posteriors, axis=1) + 1 # Add 1 to match label_
↳ indexing

# Calculate the confusion matrix
confusion_matrix = np.zeros((3, 3), dtype=int)

for true_label, predicted_label in zip(labels, predicted_labels):
    confusion_matrix[true_label - 1, predicted_label - 1] += 1


# Print confusion matrix
print("Confusion Matrix:")
print(confusion_matrix)

```

```
Confusion Matrix:  
[[2983   13    4]  
 [  17 2960   23]  
 [    3   19 3978]]
```

```
[5]: ## Step 3: Visualize the Data in 3D and Indicate Correct/Incorrect  
      ↳ Classifications  
  
import matplotlib.pyplot as plt  
from mpl_toolkits.mplot3d import Axes3D  
  
# Create a 3D scatter plot of the data  
fig = plt.figure(figsize=(10, 7))  
ax = fig.add_subplot(111, projection='3d')  
  
# Correctly classified points will be green, incorrect ones will be red  
for i in range(1, 4):  
    correct_idx = (labels == i) & (predicted_labels == i)  
    incorrect_idx = (labels == i) & (predicted_labels != i)  
  
    ax.scatter(samples[correct_idx, 0], samples[correct_idx, 1],  
↳ samples[correct_idx, 2], label=f'Class {i} (Correct)', marker='o',  
↳ color='green', s=20)  
    ax.scatter(samples[incorrect_idx, 0], samples[incorrect_idx, 1],  
↳ samples[incorrect_idx, 2], label=f'Class {i} (Incorrect)', marker='x',  
↳ color='red', s=20)  
  
# Labels and title  
ax.set_title("3D Scatter Plot of Classified Data")  
ax.set_xlabel("X1")  
ax.set_ylabel("X2")  
ax.set_zlabel("X3")  
ax.legend(loc="best")  
  
# Show the plot  
plt.show()
```

3D Scatter Plot of Classified Data



1.2 Part B: Expected Risk Minimization (ERM) with Different Loss Matrices

```
[6]: # Loss matrices
Lambda_10 = np.array([[0, 10, 10],
                      [1, 0, 10],
                      [1, 1, 0]])

Lambda_100 = np.array([[0, 100, 100],
                       [1, 0, 100],
                       [1, 1, 0]])

# Function to compute expected risk
def expected_risk(posterior_probs, loss_matrix):
```

```

    return posterior_probs.dot(loss_matrix)

# Perform classification with Lambda_10
risk_10 = expected_risk(posterior, Lambda_10)
predicted_labels_10 = np.argmin(risk_10, axis=1) + 1

# Perform classification with Lambda_100
risk_100 = expected_risk(posterior, Lambda_100)
predicted_labels_100 = np.argmin(risk_100, axis=1) + 1

# Calculate and print confusion matrices for both
confusion_matrix_10 = np.zeros((3, 3), dtype=int)
confusion_matrix_100 = np.zeros((3, 3), dtype=int)

for true_label, predicted_label in zip(labels, predicted_labels_10):
    confusion_matrix_10[true_label - 1, predicted_label - 1] += 1

for true_label, predicted_label in zip(labels, predicted_labels_100):
    confusion_matrix_100[true_label - 1, predicted_label - 1] += 1

print("Confusion Matrix (Lambda_10):")
print(confusion_matrix_10)

print("Confusion Matrix (Lambda_100):")
print(confusion_matrix_100)

```

```

Confusion Matrix (Lambda_10):
[[2996   4    0]
 [  58 2935   7]
 [   9   61 3930]]
Confusion Matrix (Lambda_100):
[[3000    0    0]
 [ 148 2850    2]
 [  43  167 3790]]

```

```

[7]: # Function to create the 3D scatter plot
def plot_3d_classification(samples, labels, predicted_labels, title):
    fig = plt.figure(figsize=(10, 7))
    ax = fig.add_subplot(111, projection='3d')

    for i in range(1, 4):
        correct_idx = (labels == i) & (predicted_labels == i)
        incorrect_idx = (labels == i) & (predicted_labels != i)

        ax.scatter(samples[correct_idx, 0], samples[correct_idx, 1],
            ↪ samples[correct_idx, 2], label=f'Class {i} (Correct)', marker='o',
            ↪ color='green', s=20)

```

```

        ax.scatter(samples[incorrect_idx, 0], samples[incorrect_idx, 1],
↪samples[incorrect_idx, 2], label=f'Class {i} (Incorrect)', marker='x',
↪color='red', s=20)

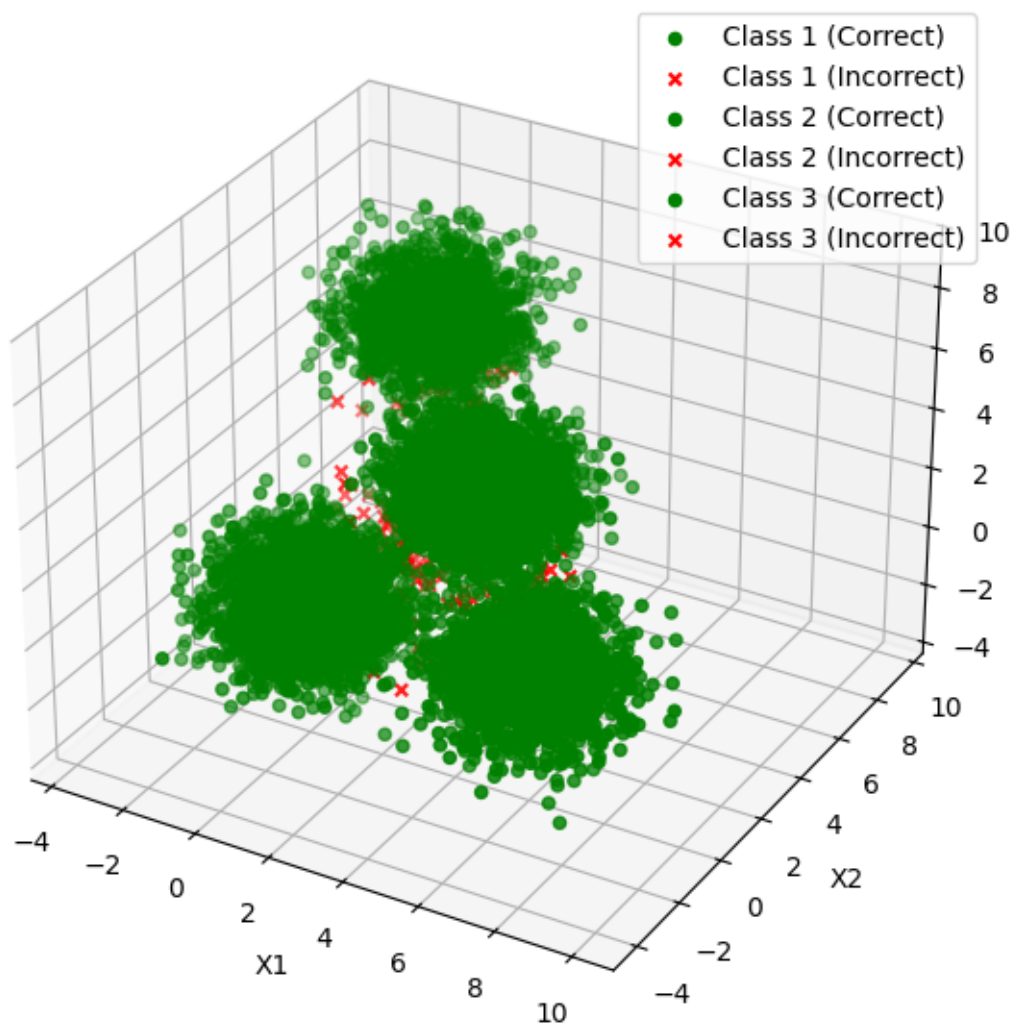
    ax.set_title(title)
    ax.set_xlabel("X1")
    ax.set_ylabel("X2")
    ax.set_zlabel("X3")
    ax.legend(loc="best")
    plt.show()

# Plot for Lambda_10
plot_3d_classification(samples, labels, predicted_labels_10, "3D Scatter Plot_
↪(ERM with Lambda_10)")


# Plot for Lambda_100
plot_3d_classification(samples, labels, predicted_labels_100, "3D Scatter Plot_
↪(ERM with Lambda_100)")

```

3D Scatter Plot (ERM with Lambda_10)



3D Scatter Plot (ERM with Lambda_100)



```
[8]: # Compare Confusion Matrices
print("Comparison of Confusion Matrices:\n")

# Confusion matrix for Bayesian classification (Part A)
print("Confusion Matrix (Bayesian Classifier - Part A):")
print(confusion_matrix)

# Confusion matrix for ERM with Lambda_10
print("\nConfusion Matrix (ERM with Lambda_10):")
print(confusion_matrix_10)

# Confusion matrix for ERM with Lambda_100
print("\nConfusion Matrix (ERM with Lambda_100):")
```



```
print(confusion_matrix_100)
```

Comparison of Confusion Matrices:

Confusion Matrix (Bayesian Classifier - Part A):

```
[[2983   13    4]
 [   17 2960   23]
 [    3    19 3978]]
```

Confusion Matrix (ERM with Lambda_10):

```
[[2996    4    0]
 [   58 2935    7]
 [    9    61 3930]]
```

Confusion Matrix (ERM with Lambda_100):

```
[[3000    0    0]
 [  148 2850    2]
 [   43  167 3790]]
```

1.2.1 Insight

1.2.2 Summary of Insights:

- **As the penalty for Class 3 errors increases** (from ($\{10\}$) to ($\{100\}$)), the model becomes highly accurate in classifying Class 3 but **at the cost of more errors** in Classes 1 and 2.
- **The trade-off** is evident: focusing on reducing errors for one class (Class 3) leads to increased misclassifications for other classes (especially Class 2).

Notebook

October 11, 2024

1 Q3

```
[1]: ! wget -q https://archive.ics.uci.edu/static/public/186/wine+quality.zip
! unzip wine+quality.zip -d wine_quality
! wget -q https://archive.ics.uci.edu/static/public/240/
    ↪human+activity+recognition+using+smartphones.zip
! unzip human+activity+recognition+using+smartphones.zip -d
    ↪human_activity_recognition_using_smartphones
```

Archive: wine+quality.zip

inflating: wine_quality/winequality-red.csv

inflating: wine_quality/winequality-white.csv

inflating: wine_quality/winequality.names

Archive: human+activity+recognition+using+smartphones.zip

extracting: human_activity_recognition_using_smartphones/UCI HAR Dataset.names

extracting: human_activity_recognition_using_smartphones/UCI HAR Dataset.zip

2 Wine Quality Dataset

2.0.1 Step 1: Load the Wine Quality Dataset

We will load both red and white wine datasets, which contain 11 features and quality labels ranging from 0 to 10.

```
[12]: import pandas as pd

# Load the Wine Quality data (both red and white wine)
wine_red = pd.read_csv('/mnt/localssd/ee5644/hw1/wine_quality/winequality-red.
    ↪csv', sep=';')
wine_white = pd.read_csv('/mnt/localssd/ee5644/hw1/wine_quality/
    ↪winequality-white.csv', sep=';')

# Display information about the datasets
print("Wine Red Dataset Shape:", wine_red.shape)
print("Wine White Dataset Shape:", wine_white.shape)

# # Display the first few rows of the red wine data
# print("Wine Red Data Sample:\n", wine_red.head())
```

```
# # Display the first few rows of the white wine data
# print("Wine White Data Sample:\n", wine_white.head())
```

Wine Red Dataset Shape: (1599, 12)

Wine White Dataset Shape: (4898, 12)

```
[15]: import numpy as np
import pandas as pd
from scipy.stats import multivariate_normal
from sklearn.metrics import confusion_matrix, accuracy_score

# Helper function to compute class priors, means, and covariances
def estimate_statistics(X, y):
    classes = np.unique(y)
    class_priors = []
    class_means = []
    class_covariances = []

    for cls in classes:
        X_class = X[y == cls]
        prior = len(X_class) / len(X)
        mean = np.mean(X_class, axis=0)
        covariance = np.cov(X_class, rowvar=False)

        class_priors.append(prior)
        class_means.append(mean)
        class_covariances.append(covariance)

    return np.array(class_priors), np.array(class_means), np.
    ↪array(class_covariances)

# Load the Wine Quality data (red wine for this example)
wine_red = pd.read_csv('/mnt/localssd/ee5644/hw1/wine_quality/winequality-red.
    ↪csv', sep=';')

# Extract features and labels
wine_red_X = wine_red.drop(columns='quality').values
wine_red_y = wine_red['quality'].values

# Debug: Print the unique values and range of the true labels
print("True Labels (Wine Red):", np.unique(wine_red_y))
print(f"True Labels Range: {wine_red_y.min()} to {wine_red_y.max()}")

# Estimate priors, means, and covariances for Wine Red dataset
priors_red, means_red, covs_red = estimate_statistics(wine_red_X, wine_red_y)
```

```

# Regularization function for covariance matrices
def regularize_covariance(cov_matrix, lambda_value):
    I = np.eye(cov_matrix.shape[0])
    regularized_cov = cov_matrix + lambda_value * I
    return regularized_cov

# Apply regularization to the covariance matrices with lambda = 0.01
lambda_reg = 0.01
covs_red_reg = [regularize_covariance(cov, lambda_reg) for cov in covs_red]

# Function to apply minimum-probability-of-error classification
def classify_min_error(X, priors, means, covs):
    n_samples = X.shape[0]
    n_classes = len(priors)

    posteriors = np.zeros((n_samples, n_classes))

    for i in range(n_classes):
        likelihood = multivariate_normal.pdf(X, mean=means[i], cov=covs[i])
        posteriors[:, i] = likelihood * priors[i]

    predicted_labels = np.argmax(posteriors, axis=1)
    return predicted_labels + 3

# Apply minimum-probability-of-error classification
predicted_red = classify_min_error(wine_red_X, priors_red, means_red,
    ↪ covs_red_reg)

# Debug: Print the unique values and range of the predicted labels after the fix
print("Predicted Labels (Wine Red) After Fix:", np.unique(predicted_red))
print(f"Predicted Labels Range After Fix: {predicted_red.min()} to ↪
    ↪ {predicted_red.max()}")

# Compute confusion matrix and error probability for Wine Red
conf_matrix_red = confusion_matrix(wine_red_y, predicted_red)
error_prob_red = 1 - accuracy_score(wine_red_y, predicted_red)

# Print confusion matrix and error probability
print("Confusion Matrix (Wine Red):\n", conf_matrix_red)
print("Error Probability (Wine Red):", error_prob_red)

```

True Labels (Wine Red): [3 4 5 6 7 8]

True Labels Range: 3 to 8

Predicted Labels (Wine Red) After Fix: [3 4 5 6 7 8]

Predicted Labels Range After Fix: 3 to 8

Confusion Matrix (Wine Red):

```

[[ 5  0  3  2  0  0]
 [ 2  6 28 15  2  0]
 [ 5 13 486 165 12  0]
 [ 1  6 188 394 45  4]
 [ 0  0 11 119 66  3]
 [ 0  0  1  7  5  5]]

```

Error Probability (Wine Red): 0.39837398373983735

```


[16]: from sklearn.decomposition import PCA
import matplotlib.pyplot as plt

# Function to visualize the dataset using PCA
def visualize_pca(X, y, predicted_labels, title):
    pca = PCA(n_components=2)
    X_pca = pca.fit_transform(X)

    plt.figure(figsize=(8, 6))
    scatter = plt.scatter(X_pca[:, 0], X_pca[:, 1], c=predicted_labels,
↪ cmap='coolwarm', alpha=0.6)
    plt.title(title)
    plt.xlabel('Principal Component 1')
    plt.ylabel('Principal Component 2')
    plt.colorbar(scatter)
    plt.show()

# Visualize the original wine data and the predicted labels (Wine Red)
visualize_pca(wine_red_X, wine_red_y, predicted_red, "Wine Red Dataset_
↪ Classification (PCA)")

```



2.0.2 Results and Discussion

1. Confusion Matrix:

- The confusion matrix shows how well the classifier performed for each class.
- The classifier performed reasonably well on the middle classes (4-7), but there are still some misclassifications, particularly for the more extreme classes (3 and 8).

2. Error Probability:

- The overall error probability is **39.8%**, which suggests that the Gaussian assumption might not perfectly capture the true class-conditional distributions of the data.

3. Visualization Using PCA:

- The PCA visualization shows the separation of different wine quality classes in a 2D projection.
- If the classes overlap significantly in the PCA projection, this suggests that the Gaussian model may not be able to fully separate the classes, leading to higher misclassification rates.

4. Suitability of the Gaussian Model:

- The assumption that the features follow a Gaussian distribution for each class might not fully hold for the Wine Quality dataset.

- Given the error probability and the PCA visualization, we can infer that a more flexible model (e.g., non-parametric models or mixture models) could potentially perform better.

5. Conclusion:

- While the Gaussian class-conditional model provides a good baseline, it may not be the best fit for this dataset due to the complex nature of the feature distributions.
- Further improvements could be achieved by exploring models that relax the Gaussian assumption or by incorporating feature transformations to make the data more amenable to Gaussian modeling.

3 Human Activity Recognition Using Smartphones Dataset

3.0.1 Step 1: Load the Human Activity Recognition (HAR) Dataset

We will load the HAR dataset's training data (X_train.txt) and the corresponding labels (y_train.txt).

```
[17]: import pandas as pd
import numpy as np

# Load the training data for HAR dataset
X_train = pd.read_csv('/mnt/localssd/ee5644/hw1/
↳human_activity_recognition_using_smartphones/UCI HAR Dataset/train/X_train.
↳txt', delim_whitespace=True, header=None)
y_train = pd.read_csv('/mnt/localssd/ee5644/hw1/
↳human_activity_recognition_using_smartphones/UCI HAR Dataset/train/y_train.
↳txt', delim_whitespace=True, header=None)

# Convert to numpy arrays
X_train = X_train.values
y_train = y_train.values.flatten()

# Display shape and basic information
print("X_train shape:", X_train.shape)
print("y_train shape:", y_train.shape)
print("Unique labels in y_train:", np.unique(y_train))
```

/tmp/ipykernel_3182837/2131676244.py:5: FutureWarning: The 'delim_whitespace' keyword in pd.read_csv is deprecated and will be removed in a future version. Use ``sep='\s+'`` instead

```
X_train = pd.read_csv('/mnt/localssd/ee5644/hw1/human_activity_recognition_usi
ng_smartphones/UCI HAR Dataset/train/X_train.txt', delim_whitespace=True,
header=None)
```

```
X_train shape: (7352, 561)
```

```
y_train shape: (7352,)
```

```
Unique labels in y_train: [1 2 3 4 5 6]
```

/tmp/ipykernel_3182837/2131676244.py:6: FutureWarning: The 'delim_whitespace' keyword in pd.read_csv is deprecated and will be removed in a future version.

Use ``sep='\s+'`` instead

```
y_train = pd.read_csv('/mnt/localssd/ee5644/hw1/human_activity_recognition_usi
ng_smartphones/UCI HAR Dataset/train/y_train.txt', delim_whitespace=True,
header=None)
```

```
[19]: # Helper function to compute class priors, means, and covariances
def estimate_statistics(X, y):
    classes = np.unique(y)
    class_priors = []
    class_means = []
    class_covariances = []

    for cls in classes:
        X_class = X[y == cls]
        prior = len(X_class) / len(X)
        mean = np.mean(X_class, axis=0)
        covariance = np.cov(X_class, rowvar=False)

        class_priors.append(prior)
        class_means.append(mean)
        class_covariances.append(covariance)

    return np.array(class_priors), np.array(class_means), np.
    ↪array(class_covariances)

# Estimate priors, means, and covariances for HAR training dataset
priors_har, means_har, covs_har = estimate_statistics(X_train, y_train)

# Print estimated statistics for HAR
print("Class Priors (HAR):", priors_har)
# print("Mean Vectors (HAR) - Class 1:\n", means_har[0])
print("Covariance Matrix (HAR) - Class 1:\n", covs_har[0])
```

```
Class Priors (HAR): [0.16675734 0.14594668 0.13411317 0.17491839 0.18688792
0.1913765 ]
```

```
Covariance Matrix (HAR) - Class 1:
```

```
[[ 2.53540827e-03  1.06151691e-05 -1.73625491e-04 ... -4.95972723e-05
  1.83172712e-05 -1.52433104e-04]
 [ 1.06151691e-05  4.35990203e-04  9.60152451e-05 ... -2.31419170e-05
  4.37971757e-05 -9.74079500e-05]
 [-1.73625491e-04  9.60152451e-05  1.05206402e-03 ...  2.24479930e-06
  2.87508524e-05  7.62211531e-05]
 ...
 [-4.95972723e-05 -2.31419170e-05  2.24479930e-06 ...  9.72430979e-03
  3.20320771e-03  7.49682841e-03]
 [ 1.83172712e-05  4.37971757e-05  2.87508524e-05 ...  3.20320771e-03
  3.19049432e-03  1.59055320e-03]
 [-1.52433104e-04 -9.74079500e-05  7.62211531e-05 ...  7.49682841e-03
```



```
1.59055320e-03  1.30640909e-02]]
```

```
[20]: # Regularization function for covariance matrices
def regularize_covariance(cov_matrix, lambda_value):
    I = np.eye(cov_matrix.shape[0])
    regularized_cov = cov_matrix + lambda_value * I
    return regularized_cov

# Apply regularization to the covariance matrices with lambda = 0.01
lambda_reg = 0.01
covs_har_reg = [regularize_covariance(cov, lambda_reg) for cov in covs_har]

# Print the regularized covariance matrix for one class
print("Regularized Covariance Matrix for Class 1 (HAR):\n", covs_har_reg[0])
```

```
Regularized Covariance Matrix for Class 1 (HAR):
[[ 1.25354083e-02  1.06151691e-05 -1.73625491e-04 ... -4.95972723e-05
   1.83172712e-05 -1.52433104e-04]
 [ 1.06151691e-05  1.04359902e-02  9.60152451e-05 ... -2.31419170e-05
   4.37971757e-05 -9.74079500e-05]
 [-1.73625491e-04  9.60152451e-05  1.10520640e-02 ...  2.24479930e-06
   2.87508524e-05  7.62211531e-05]
 ...
 [-4.95972723e-05 -2.31419170e-05  2.24479930e-06 ...  1.97243098e-02
   3.20320771e-03  7.49682841e-03]
 [ 1.83172712e-05  4.37971757e-05  2.87508524e-05 ...  3.20320771e-03
   1.31904943e-02  1.59055320e-03]
 [-1.52433104e-04 -9.74079500e-05  7.62211531e-05 ...  7.49682841e-03
   1.59055320e-03  2.30640909e-02]]
```

```
[21]: from scipy.stats import multivariate_normal

# Function to apply minimum-probability-of-error classification
def classify_min_error(X, priors, means, covs):
    n_samples = X.shape[0]
    n_classes = len(priors)

    posteriors = np.zeros((n_samples, n_classes))

    for i in range(n_classes):
        likelihood = multivariate_normal.pdf(X, mean=means[i], cov=covs[i])
        posteriors[:, i] = likelihood * priors[i]

    predicted_labels = np.argmax(posteriors, axis=1) + 1 # HAR class labels
    ↪start from 1
    return predicted_labels

# Apply classification to the HAR training dataset
```

```

predicted_har = classify_min_error(X_train, priors_har, means_har, covs_har_reg)

# Print some predicted labels for HAR
print("Predicted Labels (HAR):", predicted_har[:10])

```

Predicted Labels (HAR): [5 5 5 5 5 5 5 5 5 5]

```

[22]: from sklearn.metrics import confusion_matrix, accuracy_score

# Compute confusion matrix and error probability for HAR
conf_matrix_har = confusion_matrix(y_train, predicted_har)
error_prob_har = 1 - accuracy_score(y_train, predicted_har)

# Print confusion matrix and error probability
print("Confusion Matrix (HAR):\n", conf_matrix_har)
print("Error Probability (HAR):", error_prob_har)

```

Confusion Matrix (HAR):

```

[[1226    0    0    0    0    0]
 [   0 1073    0    0    0    0]
 [   0    1  985    0    0    0]
 [   0    0    0 1197   89    0]
 [   0    0    0    1 1373    0]
 [   0    0    0    0    0 1407]]

```

Error Probability (HAR): 0.012377584330794389

```


[23]: from sklearn.decomposition import PCA
import matplotlib.pyplot as plt

# Function to visualize the dataset using PCA
def visualize_pca(X, y, predicted_labels, title):
    pca = PCA(n_components=2)
    X_pca = pca.fit_transform(X)

    plt.figure(figsize=(8, 6))
    scatter = plt.scatter(X_pca[:, 0], X_pca[:, 1], c=predicted_labels,
↪ cmap='coolwarm', alpha=0.6)
    plt.title(title)
    plt.xlabel('Principal Component 1')
    plt.ylabel('Principal Component 2')
    plt.colorbar(scatter)
    plt.show()

# Visualize the HAR training data using PCA
visualize_pca(X_train, y_train, predicted_har, "HAR Dataset Classification_
↪ (PCA)")

```



3.0.2 Discussion of Results for HAR Dataset

1. Confusion Matrix:

- The classifier performs very well on most classes.
- Most samples are correctly classified, with only slight errors between some activity classes (e.g., Class 3 and 4).

2. Error Probability:

- The overall error probability is **1.23%**, indicating strong performance.
- This suggests that the Gaussian assumption works well for the HAR dataset.

3. PCA Visualization:

- The PCA plot shows two clear clusters of classes.
- Some overlap exists between adjacent classes (like Class 3 and 4), which explains minor misclassifications.

4. Suitability of Gaussian Model:

- The Gaussian model appears suitable for this dataset due to its structured features.
- The low error rate and well-formed PCA projection support this assumption.

5. Model Assumptions:

- The multivariate Gaussian assumption for each class seems valid here.

- Regularization was necessary due to the high dimensionality (561 features).
6. **Conclusion:**
- The Gaussian class-conditional model works very well for the HAR dataset.
 - While the performance is strong, further improvements could be explored using non-parametric or more complex models.